

Faculty of Information Technology,
Monash University

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

This material has been reproduced and communicated to you by or on behalf of Monash University pursuant to Part VB of the Copyright Act 1968 (the Act). The material in this communication may be subject to copyright under the Act. Any further reproduction or communication of this material by you may be the subject of copyright protection under the Act. Do not remove this notice

FIT2004: Algorithms and Data Structures

Week 7: Dynamic Programming Graph Algorithms

Outline

Divide and conquer
(W 1-3)

Greedy algorithms
(W 4-5)

Dynamic programming
(W 6-7)

Network flow
(W 8-9)

Data structures
(W 10-11)

- Last Lecture: DP algorithms
 - Coins Change
 - Knapsack
 - Edit Distance
- Today's Lecture: **DP graph algorithms**
 - Shortest path in graphs with negative weights
 - All-pairs shortest paths
 - Transitive Closure

FIT2004, Lecture 7 - DP Graph Algorithms


Outline

1. Shortest path in a graph with negative weights (Bellman-Ford Algorithm)
2. All-pairs shortest paths (Floyd-Warshall Algorithm)
3. Transitive Closure

FIT2004, Lecture 7 - DP Graph Algorithms

Shortest path (negative weights)


- What is the shortest distance from s to v in this graph?
- If Dijkstra's algorithm is used on this graph, what will it output as being the shortest path from s to v?
- Dijkstra's algorithm is **not guaranteed** to output the correct answer when there are negative weights.



FIT2004, Lecture 7 - DP Graph Algorithms

Shortest path (negative weights)


- What is the shortest distance from s to v in this graph?
- If Dijkstra's algorithm is used on this graph, what will it output as being the shortest path from s to v?
- Dijkstra's algorithm is **not guaranteed** to output the correct answer when there are negative weights.



FIT2004, Lecture 7 - DP Graph Algorithms

Shortest path (negative weights)


- What is the shortest distance from s to v in this graph?
- If Dijkstra's algorithm is used on this graph, what will it output as being the shortest path from s to v?
- Dijkstra's algorithm is **not guaranteed** to output the correct answer when there are negative weights.



FIT2004, Lecture 7 - DP Graph Algorithms

Shortest path (negative weights)

What is the shortest distance from s to x in this graph?




FIT2004, Lecture 7 - DP Graph Algorithms

Shortest path (negative weights)

What is the shortest distance from s to x in this graph?

- Not well-defined:
 - From s, it is possible to reach the negative cycle $u \rightarrow v \rightarrow t$, and from this cycle it is possible to reach x.
 - Given any path P, it is possible to obtain an alternative path P' with smaller total weight than P: P' goes from s to the negative cycle, include **as many repetitions of the negative cycle as necessary**, and then reaches x from the negative cycle.



FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm

- Bellman-Ford's algorithm **works correctly even with negative edge weights (as long as there are no negative weight cycles)**.
- Bellman-Ford algorithm returns:
 - shortest distances from s to all vertices in the graph if there are no negative cycles that are reachable from s.
 - **an error if there is a negative cycle reachable from s** (i.e., can be used to detect negative cycles).
- It can be modified to return all valid shortest distances, and minus ∞ for vertices which are affected by the negative cycle.

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Core Idea

- Idea: If no negative cycles are reachable from node s , then for every node t that is reachable from s , there is a shortest path from s to t that is **simple** (i.e., no nodes are repeated).
- Can the shortest path from s to t have a positive cycle?

No, the **shortest path** from s to t will never include a cycle (either positive or negative) if the goal is to minimize the total cost of the path.


Why is that?

- A cycle should revisit at least one node.
- If the cycle is positive (i.e., the total weight of the cycle is > 0), then going through it increases the total cost of the path.
- So a shorter path can be found by just skipping that cycle.
- Therefore, the shortest path will always be simple (i.e., no repeated nodes).

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Core Idea

- If no negative cycles are reachable from node s , then for every node t that is reachable from s there is a shortest path from s to t that is **simple** (i.e., no nodes are repeated).
- Cycles with positive weight cannot be a part of a shortest path.
- Given a shortest path that contains cycles of weight 0, the cycles can be removed to obtain an alternative shortest path that is simple.



- Note that any **simple** path has at most $V-1$ edges.

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm

A fact from Week 5: If P is a shortest path from s to u , and v is the last vertex on P before u , then the part of P from s to v is also a shortest path.

Proof:


Suppose there was a shorter path from s to v , say Q .

If it's the shortest to reach u ,

$$\text{weight}(Q) + w(v, u) < \text{weight}(P)$$

But P is the shortest path from s to u . (Contradiction)

Therefore, the part of P from s to v is also a shortest path.



FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm

- Bellman-Ford was one of the first applications of dynamic programming.



Richard Bellman



Lester Ford Jr.

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm

- For a source node s , let $OPT(i,v)$ denote the minimum weight of a $s \rightarrow v$ path with at most i edges. (Here v is for vertices)
- Let P be an optimal path with at most i edges that achieves total weight $OPT(i,v)$:
 - If P has at most $i-1$ edges, then $OPT(i,v) = OPT(i-1,v)$.
 - If P has exactly i edges and (u,v) is the last edge of P , then $OPT(i,v) = OPT(i-1,u) + w(u,v)$, where $w(u,v)$ denotes the weight of edge (u,v) .
- Recursive formula for dynamic programming:

$$OPT(i,v) = \min(OPT(i-1,v), \min_{u:(u,v) \in E} (OPT(i-1,u) + w(u,v)))$$

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm

Uses array $M[0 \dots V-1, 1 \dots V]$

Initialize $M[0,s] = 0$, for all other vertices $M[0,v] = \text{infinity}$

for $i = 1$ to $V-1$:

for each vertex v :

 Compute $M[i,v]$ using the recurrence

return $M[V-1, 1 \dots V]$

$$OPT(i,v) = \min(OPT(i-1,v), \min_{u:(u,v) \in E} (OPT(i-1,u) + w(u,v)))$$

What is the time complexity of Bellman-ford algorithm?

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm

Uses array $M[0 \dots V-1, 1 \dots V]$

Initialize $M[0, s] = 0$, for all other vertices $M[0, v] = \text{infinity}$

for $i = 1$ to $V-1$:

for each vertex v :

 Compute $M[i, v]$ using the recurrence

return $M[V-1, 1 \dots V]$

What is the time complexity of Bellman-ford algorithm?

Time Complexity:

$O(VE)$

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm

- Commonly, a more space-efficient version of Bellman-Ford algorithm is implemented.
- $V-1$ iterations are performed, but the value i is used just as a counter, and in each iteration, for each node v , we use the rule

$$M[v] = \min(M[v], \min_{u: (u,v) \in E} (M[u] + w(u, v)))$$

- In some cases, this version also provides a speed-up (but no improvement in the worst-case time complexity).

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm


- $V-1$ iterations are performed, but the value i is used just as a counter, and in each iteration, for each node v , we use following update rule for the distance:

$$dist[v] = \min(dist[v], \min_{u:(u,v) \in E} (dist[u] + w(u,v)))$$

- If vertices are updated in the order s, v, x, t, u , then we are done after 1 iteration.
- On the other hand, if vertices are updated in the order u, t, x, v, s , then we need 4 iterations to get the right result.
- We will analyse the early stopping condition later on.

i	s	v	x	t	u
1	0	7	16	23	21
2					
3					
4					

u	t	x	v	s
∞	∞	∞	7	0
∞	∞	16	7	0
∞	23	16	7	0
21	23	16	7	0



FIT2004, Lecture 7 - DP Graph Algorithms


Bellman-Ford Algorithm: Example

Initialize:

- For each vertex a in the graph
 - $dist(s,a) = \infty$
- $dist(s,s) = 0$

Consider the following operation (relaxation):

- For each edge (a, b) in the graph
 - $dist(s, b) = \min(dist(s,b), dist(s,a) + w(a,b))$



Assume the following order:


$(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)$

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

First iteration:

Relaxing incoming edges of node u



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

First iteration:

Done relaxing incoming edges of node u



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

First iteration:

Relaxing incoming edges of node v



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

First iteration:

Done relaxing incoming edges of node v



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

First iteration:

Relaxing incoming edges of node t



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

First iteration:

Done relaxing incoming edges of node t



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

First iteration:

Relaxing incoming edges of node x



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), **(u,x), (v,x)**

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

First iteration:

Done relaxing incoming edges of node x




Assume the following order:

(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

First iteration finished:



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

Second iteration:

Relaxing incoming edges of node u



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

Second iteration:

Done relaxing incoming edges of node u



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

Second iteration:

Relaxing incoming edges of node v



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

Second iteration:

Done relaxing incoming edges of node v



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

Second iteration:

Relaxing incoming edges of node t



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

Second iteration:

Done Relaxing incoming edges of node t



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

Second iteration:

Relaxing incoming edges of node x



Assume the following order:


(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

Second iteration:

Done Relaxing incoming edges of node x




Assume the following order:

(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

Second iteration finished:



Assume the following order:

(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)


FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm: Example

Third iteration:

Speeding things up: All edges relaxation in the third iteration do not change anything.

Early Stop Condition: If nothing changes in one iteration, it is possible to stop the execution of the Bellman-Ford algorithm and output the current values.




Assume the following order:

(s,u), (t,u), (s,v), (u,v), (u,t), (v,t), (x,t), (u,x), (v,x)

FIT2004, Lecture 7 - DP Graph Algorithms

Finding Negative Cycles

- If V-th iteration reduces the distance of a vertex, it tells us that there is a shorter path with at least V edges.
 - This implies that there is a negative cycle.
- For example, consider the graph with 4 vertices s, u, v, and t and assume we have run (V-1 = 3) iterations.
- In the 4th iteration, the weight of at least one vertex will be reduced (due to the presence of a negative cycle).
- **Important: Bellman-Ford Algorithm finds negative cycles only if such cycle is reachable from the source vertex.**
 - E.g., if x is the source vertex, the algorithm will not detect the negative cycle




FIT2004, Lecture 7 - DP Graph Algorithms

Detecting Negative Cycles

To detect if a graph G has a negative cycle:

Just add one extra node (z) to G and edges from it to every other node, and run Bellman-Ford on the added node, z .



FIT2004, Lecture 7 - DP Graph Algorithms

Updated Bellman-Ford Algorithm

```
# STEP 1: Initializations
dist[1...V] = infinity
pred[1...V] = Null
dist[s] = 0
# STEP 2: Iteratively estimate dist[v] (from source s)
for i = 1 to V-1:
    for each edge <u,v> in the whole graph:
        est = dist[u] + w(u,v)
        if est < dist[v]:
            dist[v] = est
            pred[v] = u

# STEP 3: Checks and returns false if a negative weight cycle
# is along the path from s to any other vertex
for each edge <u,v> in the whole graph:
    if dist[u] + w(u,v) < dist[v]:
        return error; # negative edge cycle found in this graph

return dist[...], pred[...]
```


Time Complexity:

$O(VE)$

FIT2004, Lecture 7 - DP Graph Algorithms

Bellman-Ford Algorithm

- For this space-efficient version of Bellman-Ford algorithm, there is a guarantee that **after i iterations $\text{dist}[v]$ is no larger than the total weight of the shortest path from s to v that uses at most i edges.**
- But there is no guarantee that these two values are equal after i iterations: depending on the order in which the edges are relaxed, the path P from s to v that has weight $\text{dist}[v]$ could already contain more than i edges after the i -th iteration of the outer loop.
 - e.g., in the graph that we followed a detailed execution of Bellman-Ford, the path from s to t already has two edges after just one iteration.



FIT2004, Lecture 7 - DP Graph Algorithms

Handling Negative Cycles

- How could we modify Bellman-Ford to determine **which** vertices have valid distances, and which are affected by the negative cycle?
- Execute V additional iterations, and for each node whose distance would be updated, just mark its distance as $-\infty$.
 - By continuing to relax the graph for V more iterations, the effect of negative cycles is allowed to propagate across the graph.
 - Therefore, any vertex that can be reached from a negative cycle will eventually get $-\infty$ as its distance.

FIT2004, Lecture 7 - DP Graph Algorithms

Outline

1. Shortest path in a graph with negative weights (Bellman-Ford Algorithm)
2. All-pairs shortest paths (Floyd-Warshall Algorithm)
3. Transitive Closure

FIT2004, Lecture 7 - DP Graph Algorithms

All-Pairs Shortest Paths

Problem

- Return shortest distances between **all** pairs of vertices in a connected graph.

For unweighted graphs:

- For each vertex v in the graph
 - Call Breadth-First Search (BFS) for v

Time complexity:

$O(V(V+E)) = O(V^2 + EV) \rightarrow O(EV)$ [for connected graphs, $O(V) \leq O(E)$]

For dense graphs: $E \approx O(V^2)$, therefore total cost is $O(V^3)$ for dense graphs

FIT2004, Lecture 7 - DP Graph Algorithms

All-Pairs Shortest Paths

For weighted graphs (with non-negative weights):

- For each vertex v in the graph
 - Call Dijkstra's algorithm for v

Time complexity:

$$O(V(E \log V)) = O(EV \log V)$$

For dense graphs: $O(V^3 \log V)$ since $E \approx O(V^2)$

FIT2004, Lecture 7 - DP Graph Algorithms

All-Pairs Shortest Paths

For weighted graphs (allowing negative weights):

- For each vertex v in the graph
 - Call Bellman-Ford algorithm for v

Time complexity:

$$O(V(VE)) = O(V^2 E)$$

For dense graphs: $O(V^4)$ since $E \approx O(V^2)$

Can we do better?

- Yes, **Floyd-Warshall algorithm** returns all-pairs shortest distances in $O(V^3)$ (even for graphs with negative weights).

FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm

- Algorithm based on dynamic programming.



Robert W. Floyd
(Turing Award 1978)



Stephen Warshall

FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm

- Algorithm based on dynamic programming.
- Computes the shortest distance between every pair of nodes as long as there are no negative weight cycles.
- If the graph has a negative cycle, it will always be detected.
 - Is this similar to Bellman-Ford algorithm?
 - ✖ No.
 - ✖ Bellman-Ford **only** detects negative cycles that are reachable from the source node.
- For a graph without negative cycles, after the **k-th** iteration, **dist[i][j]** contains the weight of the shortest path from node i to node j that only uses intermediate nodes from the set **{1,..., k}**.

FIT2004, Lecture 7 - DP Graph Algorithms

Key idea behind Floyd-Warshall Algorithm

- For each pair of vertices (i, j) , find whether there is a shorter path from i to j by going through an intermediate vertex k .
- All variables i, j, k refer to vertices.
- At each step, it updates:

$$\text{dist}[i][j] = \min(\text{dist}[i][j], \text{dist}[i][k] + \text{dist}[k][j])$$

- It does this for every possible value of k from 0 to $V - 1$.
- After the algorithm finishes, check for all vertices:

$\text{dist}[i][i] < 0$ which tells that a Negative weight cycle is detected.

FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm

- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

Assume that the outer for-loop will access vertices in the order A, B, C, D

First iteration of outer loop (i.e., k is A): **A to A**

	To vertex			
	A	B	C	D
From vertex	A	0	Inf	-2
	B	4	0	3
	C	Inf	Inf	0
	D	Inf	-1	0

$k =$ A B C D


$i =$ A B C D

$j =$ A B C D

$\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$

$\text{dist}[A][A] > \text{dist}[A][A] + \text{dist}[A][A]$

$0 > 0+0 \rightarrow$ no update



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm


- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✦ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

The outer for-loop will access vertices in the order A, B, C, D

First iteration of outer loop (i.e., k is A): **A to B**

	A	B	C	D
A	0	Inf	-2	Inf
B	4	0	3	Inf
C	Inf	Inf	0	2
D	Inf	-1	Inf	0

$k =$ A B C D
 $i =$ A B C D
 $j =$ A B C D
 $\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$
 $\text{dist}[A][B] > \text{dist}[A][A] + \text{dist}[A][B]$
 $\infty > 0 + \infty \rightarrow \text{no update}$



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm


- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✦ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

The outer for-loop will access vertices in the order A, B, C, D

First iteration of outer loop (i.e., k is A): **A to C**

	A	B	C	D
A	0	Inf	-2	Inf
B	4	0	3	Inf
C	Inf	Inf	0	2
D	Inf	-1	Inf	0

$k =$ A B C D
 $i =$ A B C D
 $j =$ A B C D
 $\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$
 $\text{dist}[A][C] > \text{dist}[A][A] + \text{dist}[A][C]$
 $-2 > 0 - 2 \rightarrow \text{no update}$



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm


- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✦ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

The outer for-loop will access vertices in the order A, B, C, D

First iteration of outer loop (i.e., k is A): **A to D**

	A	B	C	D
A	0	Inf	-2	Inf
B	4	0	3	Inf
C	Inf	Inf	0	2
D	Inf	-1	Inf	0

$k =$ A B C D
 $i =$ A B C D
 $j =$ A B C **D**
 $\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$
 $\text{dist}[A][D] > \text{dist}[A][A] + \text{dist}[A][D]$
 $\infty > 0 + \infty \rightarrow \text{no update}$



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm


- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✦ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

The outer for-loop will access vertices in the order A, B, C, D

First iteration of outer loop (i.e., k is A): **B to A**

	A	B	C	D
A	0	Inf	-2	Inf
B	4	0	3	Inf
C	Inf	Inf	0	2
D	Inf	-1	Inf	0

$k =$ A B C D
 $i =$ A **B** C D
 $j =$ A B C D
 $\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$
 $\text{dist}[B][A] > \text{dist}[B][A] + \text{dist}[A][A]$
 $4 > 4 + 0 \rightarrow \text{no update}$



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm


- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✦ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

The outer for-loop will access vertices in the order A, B, C, D

First iteration of outer loop (i.e., k is A): **B to C**

	A	B	C	D
A	0	Inf	-2	Inf
B	4	0	3	Inf
C	Inf	Inf	0	2
D	Inf	-1	Inf	0

$k =$ A B C D
 $i =$ A B C D Let's speed up
 $j =$ A B C D
 $\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$
 $\text{dist}[B][C] > \text{dist}[B][A] + \text{dist}[A][C]$
 $3 > 4 - 2 \rightarrow \text{yes, then update}$



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm


- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✦ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

The outer for-loop will access vertices in the order A, B, C, D

First iteration of outer loop (i.e., k is A): **B to C**

	A	B	C	D
A	0	Inf	-2	Inf
B	4	0	3	Inf
C	Inf	Inf	0	2
D	Inf	-1	Inf	0

$k =$ A B C D
 $i =$ A B C D Let's speed up
 $j =$ A B C D
 $\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$
 $\text{dist}[B][C] > \text{dist}[B][A] + \text{dist}[A][C]$
 $3 > 4 - 2 \rightarrow \text{yes, then update}$



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm


- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✦ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

The outer for-loop will access vertices in the order A, B, C, D

First iteration of outer loop (i.e., k is A): **B to C**

	A	B	C	D
A	0	Inf	-2	Inf
B	4	0	2	Inf
C	Inf	Inf	0	2
D	Inf	-1	Inf	0

$k =$ A B C D
 $i =$ A B C D Let's speed up
 $j =$ A B C D
 $\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$
 $\text{dist}[B][C] > \text{dist}[B][A] + \text{dist}[A][C]$
 $3 > 4-2 \rightarrow$ yes, then update



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm


- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✦ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

The outer for-loop will access vertices in the order A, B, C, D

First iteration of outer loop (i.e., k is A): **B to D**

	A	B	C	D
A	0	Inf	-2	Inf
B	4	0	2	Inf
C	Inf	Inf	0	2
D	Inf	-1	Inf	0

$k =$ A B C D
 $i =$ A B C D Let's speed up
 $j =$ A B C D
 $\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$
 $\text{dist}[B][D] > \text{dist}[B][A] + \text{dist}[A][D]$
 $\infty > 4+\infty \rightarrow$ no update



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm

- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✦ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

First iteration of outer loop (i.e., k is A):


The remaining pairs are (C,A), (C,B), (C,C), (C,D), (D,A), (D,B), (D,C) and (D,D).

It is not possible to improve the distance between any of these pairs of nodes using only A as intermediate.

Hence, we move to next iteration.

	A	B	C	D
A	0	Inf	-2	Inf
B	4	0	2	Inf
C	Inf	Inf	0	2
D	Inf	-1	Inf	0

$k =$ A B C D
 $i =$ A B C D
 $j =$ A B C D
 $\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm


- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✦ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

The outer for-loop will access vertices in the order A, B, C, D

Second iteration of outer loop (i.e., k is B): D to A and D to C

	A	B	C	D
A	0	Inf	-2	Inf
B	4	0	2	Inf
C	Inf	Inf	0	2
D	Inf	-1	Inf	0

$k =$ A B C D
 $i =$ A B C D
 $j =$ A B C D
 $\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$
 $\text{dist}[D][A] > \text{dist}[D][B] + \text{dist}[B][A] \rightarrow \infty > -1 + 4$, update
 $\text{dist}[D][C] > \text{dist}[D][B] + \text{dist}[B][C] \rightarrow \infty > -1 + 2$, update



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm

- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✱ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

The outer for-loop will access vertices in the order A, B, C, D

Second iteration of outer loop (i.e., k is B): D to A and D to C

	A	B	C	D
A	0	Inf	-2	Inf
B	4	0	2	Inf
C	Inf	Inf	0	2
D	3	-1	1	0

$k =$ A B C D


$i =$ A B C D

$j =$ A B C D

$\text{dist}[i][j] > \text{dist}[i][k] + \text{dist}[k][j]$

$\text{dist}[D][A] > \text{dist}[D][B] + \text{dist}[B][A] \rightarrow \infty > -1 + 4 = 3$

$\text{dist}[D][C] > \text{dist}[D][B] + \text{dist}[B][C] \rightarrow \infty > -1 + 2 = 1$



FIT2004, Lecture 7 - DP Graph Algorithms


Floyd-Warshall Algorithm

- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✱ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

Assume that the outer for-loop will access vertices in the order A, B, C, D

Using nodes from {A, B, C} as intermediates, it is possible to update the following distances:

	A	B	C	D
A	0	Inf	-2	0
B	4	0	2	4
C	Inf	Inf	0	2
D	3	-1	1	0



FIT2004, Lecture 7 - DP Graph Algorithms


Floyd-Warshall Algorithm

- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✦ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

Assume that the outer for-loop will access vertices in the order A, B, C, D

Using nodes from {A, B, C, D} as intermediates, it is possible to update the following distances:

	A	B	C	D
A	0	-1	-2	0
B	4	0	2	4
C	5	1	0	2
D	3	-1	1	0



FIT2004, Lecture 7 - DP Graph Algorithms


Floyd-Warshall Algorithm

- Initialise adjacency matrix called $\text{dist}[][]$ considering adjacent edges only
- For each vertex k in the graph
 - For each pair of vertices i and j in the graph
 - ✦ If $\text{dist}[i][k] + \text{dist}[k][j] < \text{dist}[i][j]$ // i.e., $\text{dist}(i \rightarrow k \rightarrow j)$ is smaller than the current $\text{dist}(i \rightarrow j)$
 - Update $\text{dist}[i][j] = \text{dist}[i][k] + \text{dist}[k][j]$ // create shortcut $i \rightarrow j$ with weight equal to $\text{dist}(i \rightarrow k \rightarrow j)$

Assume that the outer for-loop will access vertices in the order A, B, C, D

Final Solution:

	A	B	C	D
A	0	-1	-2	0
B	4	0	2	4
C	5	1	0	2
D	3	-1	1	0



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm

```

dist[][] = E # Initialize adjacency matrix using E
for vertex k in 1..V:
    #Invariant: dist[i][j] corresponds to the shortest path from i
    to j considering the intermediate vertices 1 to k-1
    for vertex i in 1..V:
        for vertex j in 1..V:
            dist[i][j] = min(dist[i][j], dist[i][k] + dist[k][j])

```

Time Complexity:

$O(V^3)$

Space Complexity:

$O(V^2)$

FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm: Correctness

Invariant: $\text{dist}[i][j]$ corresponds to the shortest path from i to j considering only intermediate vertices 1 to $k-1$.

Base Case $k = 1$ (i.e. there are no intermediate vertices yet):

- It is true because $\text{dist}[][]$ is initialized based only on the adjacent edges.

Inductive Step:

- Assume $\text{dist}[i][j]$ is the shortest path from i to j detouring through only vertices 1 to $k-1$.
- Adding the k -th vertex to the “detour pool” can only help if the best path detours through k .
- Thus, minimum of $\text{dist}(i \rightarrow k \rightarrow j)$ and $\text{dist}(i \rightarrow j)$ gives the minimum distance from i to j considering the intermediate vertices 1 to k .

FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm: Correctness

Invariant: $\text{dist}[i][j]$ corresponds to the shortest path from i to j considering only intermediate vertices 1 to $k-1$


- Adding the k -th vertex to the “detour pool” can only help if the best path detours through k .
- We already know the best way to get from i to k (using only vertices in $1\dots k-1$) and we know the best way to get from k to j (using only vertices in $1\dots k-1$).
- Thus, minimum of $\text{dist}(i \rightarrow k \rightarrow j)$ and $\text{dist}(i \rightarrow j)$ gives the minimum distance from i to j considering the intermediate vertices 1 to k .

FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm: Negative Cycles

- If there is a negative cycle, there will be a vertex v such that $\text{dist}[v][v]$ is negative.
- Look at the diagonal of the final matrix and **return error** if a negative value is found.
- How could you modify the algorithm to return the **paths**?
 - Add a “predecessor” matrix that stores the path information when the shortest distances are updated.

	A	B	C	D
A	0	-5	-3	-5
B	5	-1	1	-1
C	3	-3	-1	-3
D	4	-2	0	-2



FIT2004, Lecture 7 - DP Graph Algorithms


Outline

1. Shortest path in a graph with negative weights
2. All-pairs shortest paths
3. **Transitive Closure**

FIT2004, Lecture 7 - DP Graph Algorithms

Transitive Closure of a Graph

- Given a graph $G = (V, E)$, its transitive closure is another graph (V, E') that contains the same vertices V but contains an edge from node u to node v if there is a path from u to v in the original graph.
- **Solution:** Assign each edge a weight 1 and then apply Floyd-Warshall algorithm. If $\text{dist}[i][j]$ is not infinity, this means that there is a path from i to j in the original graph. (Or just maintain True and False as shown next).



FIT2004, Lecture 7 - DP Graph Algorithms

Floyd-Warshall Algorithm for Transitive Closure

```
# Modify Floyd-Warshall Algorithm to compute Transitive Closure
# initialization
for vertex i in 1..V:
    for vertex j in 1..V:
        if there is an edge between i and j or i == j:
            TC[i][j] = True
        else:
            TC[i][j] = False
for vertex k in 1..V:
    # Invariant: TC[i][j] corresponds to the existence of path from i to j considering the intermediate
    # vertices 1 to k-1
    for vertex i in 1..V:
        for vertex j in 1..V:
            TC[i][j] = TC[i][j] or (TC[i][k] and TC[k][j])
```

Time Complexity:

$O(V^3)$

Space Complexity:

$O(V^2)$

FIT2004, Lecture 7 - DP Graph Algorithms

Reading

- Course Notes: Chapter 8
- You can also check algorithms' textbooks for contents related to this lecture, e.g.:
 - CLRS: Sections 24.1 and 25.2
 - KT: Sections 6.8, 6.9 and 6.10
 - Rou: Chapter 18

FIT2004, Lecture 7 - DP Graph Algorithms

Concluding Remarks

Take home message

- Dijkstra's algorithm works only for graphs with non-negative weights.
- Bellman-Ford computes shortest paths in graphs with negative weights in $O(VE)$ and can also detect the negative cycles that are reachable.
- Floyd-Warshall Algorithm computes all-pairs shortest paths and transitive closure in $O(V^3)$.

Things to do (this list is not exhaustive)

- Go through recommended reading and make sure you understand why the algorithms are correct.
- Implement Bellman-Ford and Floyd-Warshall Algorithms.

Coming Up Next

- Network Flow